

---

### Task 1: Vehicle Maintenance Data Ingestion

- Use the following CSV data representing vehicle maintenance records:

```
VehicleID,Date,ServiceType,ServiceCost,Mileage
V001,2024-04-01,Oil Change,50.00,15000
V002,2024-04-05,Tire Replacement,400.00,30000
V003,2024-04-10,Battery Replacement,120.00,25000
V004,2024-04-15,Brake Inspection,200.00,40000
V005,2024-04-20,Oil Change,50.00,18000
```

- Ingest this CSV data into a Delta table in Databricks.
- Add error handling for cases where the file is missing or contains incorrect data, and log any such issues.

### Task 2: Data Cleaning

- Clean the vehicle maintenance data:
  - Ensure that the `ServiceCost` and `Mileage` columns contain valid positive values.
  - Remove any duplicate records based on `VehicleID` and `Date`.
  - Save the cleaned data to a new Delta table.

### Task 3: Vehicle Maintenance Analysis

- Create a notebook to analyze the vehicle maintenance data:
  - Calculate the total maintenance cost for each vehicle.
  - Identify vehicles that have exceeded a certain mileage threshold (e.g., 30,000 miles) and might need additional services.
  - Save the analysis results to a Delta table.

### Task 5: Data Governance with Delta Lake

- Enable Delta Lake's data governance features:
  - Use `VACUUM` to clean up old data from the Delta table.
  - Use `DESCRIBE HISTORY` to check the history of updates to the maintenance records.

---

### Task 1: Movie Ratings Data Ingestion

- Use the following CSV data to represent movie ratings by users:

```
UserID,MovieID,Rating,Timestamp
U001,M001,4,2024-05-01 14:30:00
U002,M002,5,2024-05-01 16:00:00
U003,M001,3,2024-05-02 10:15:00
U001,M003,2,2024-05-02 13:45:00
U004,M002,4,2024-05-03 18:30:00
```

- Ingest this CSV data into a Delta table in Databricks.
- Ensure proper error handling for missing or inconsistent data, and log errors accordingly.

### Task 2: Data Cleaning

- Clean the movie ratings data:
  - Ensure that the `Rating` column contains values between 1 and 5.
  - Remove any duplicate entries (same `UserID` and `MovieID`).
  - Save the cleaned data to a new Delta table.

### Task 3: Movie Rating Analysis

- Create a notebook to analyze the movie ratings:
  - Calculate the average rating for each movie.
  - Identify the movies with the highest and lowest average ratings.
  - Save the analysis results to a Delta table.

### Task 4: Time Travel and Delta Lake History

- Implement Delta Lake's time travel feature:
  - Perform an update to the movie ratings data (e.g., change a few ratings).
  - Roll back to a previous version of the Delta table to retrieve the original ratings.
  - Use `DESCRIBE HISTORY` to view the history of changes to the Delta table.

### Task 5: Optimize Delta Table

- Apply optimizations to the Delta table:
  - Implement `Z-ordering` on the `MovieID` column to improve query performance.
  - Use the `OPTIMIZE` command to compact the data and improve performance.
  - Use `VACUUM` to clean up older versions of the table.

---

---

## Task 1: Data Ingestion - Reading Data from Various Formats

1. **Ingest data from different formats** (CSV, JSON, Parquet, Delta table):
  - **CSV Data:** Use the following CSV data to represent student information:

```
StudentID,Name,Class,Score
S001,Anil Kumar,10,85
S002,Neha Sharma,12,92
S003,Rajesh Gupta,11,78
```

- **JSON Data:** Use the following JSON data to represent city information:

```
[
  {"CityID": "C001", "CityName": "Mumbai", "Population": 20411000},
  {"CityID": "C002", "CityName": "Delhi", "Population": 16787941},
  {"CityID": "C003", "CityName": "Bangalore", "Population": 8443675}
]
```

- **Parquet Data:** Use a dataset containing data about hospitals stored in Parquet format. Write code to load this data into a DataFrame.
- **Delta Table:** Load a Delta table containing hospital records, ensuring you include proper error handling in case the table does not exist.

## Task 2: Writing Data to Various Formats

1. **Write data from the following DataFrames to different formats:**

- **CSV:** Write the student data (from Task 1) to a CSV file.
- **JSON:** Write the city data (from Task 1) to a JSON file.
- **Parquet:** Write the hospital data (from Task 1) to a Parquet file.
- **Delta Table:** Write the hospital data to a Delta table.

### Task 3: Running One Notebook from Another

#### 1. Create two notebooks:

- Notebook A: Ingest data from a CSV file, clean the data (remove duplicates, handle missing values), and save it as a Delta table.
- Notebook B: Perform analysis on the Delta table created in Notebook A (e.g., calculate the average score of students) and write the results to a new Delta table.

#### 2. Run Notebook B from Notebook A:

- Implement the logic to call and run Notebook B from within Notebook A.

### Task 4: Databricks Ingestion

#### 1. Read data from the following sources:

- CSV file from Azure Data Lake.
- JSON file stored on Databricks FileStore.
- Parquet file from an external data source (e.g., AWS S3).
- Delta table stored in a Databricks-managed database.

#### 2. Write the cleaned data to each of the formats listed above (CSV, JSON, Parquet, and Delta) after performing some basic transformations (e.g., filtering rows, calculating totals).

---

### Additional Tasks:

- **Optimization Task:** Once the data is written to a Delta table, optimize it using Delta Lake's `OPTIMIZE` command.
  - **Z-ordering Task:** Apply Z-ordering on the `CityName` or `Class` columns for faster querying.
  - **Vacuum Task:** Use the `VACUUM` command to clean up old versions of the Delta table.
- 

### Exercise 1: Creating a Complete ETL Pipeline using Delta Live Tables (DLT)

#### Objective:

Learn how to create an end-to-end ETL pipeline using Delta Live Tables.

#### Tasks:

##### 1. Create Delta Live Table (DLT) Pipeline:

- Set up a DLT pipeline for processing transactional data. Use sample data representing daily customer transactions.

```
TransactionID,TransactionDate,CustomerID,Product,Quantity,Price
1,2024-09-01,C001,Laptop,1,1200
2,2024-09-02,C002,Tablet,2,300
3,2024-09-03,C001,Headphones,5,50
4,2024-09-04,C003,Smartphone,1,800
5,2024-09-05,C004,Smartwatch,3,200
```

- Define the pipeline steps:
  - **Step 1:** Ingest raw data from CSV files.
  - **Step 2:** Apply transformations (e.g., calculate total transaction amount).
  - **Step 3:** Write the final data into a Delta table.

## 2. Write DLT in Python:

- Implement the pipeline using **DLT in Python**. Define the following tables:
  - **Raw Transactions Table:** Read data from the CSV file.
  - **Transformed Transactions Table:** Apply transformations (e.g., calculate total amount: `Quantity * Price`).

## 3. Write DLT in SQL:

- Implement the same pipeline using **DLT in SQL**. Use SQL syntax to define tables, transformations, and outputs.

## 4. Monitor the Pipeline:

- Use Databricks' DLT UI to monitor the pipeline and check the status of each step.

---

## Exercise 2: Delta Lake Operations - Read, Write, Update, Delete, Merge

### Objective:

Work with Delta Lake to perform read, write, update, delete, and merge operations using both PySpark and SQL.

### Tasks:

#### 1. Read Data from Delta Lake:

- Read the transactional data from the Delta table you created in the first exercise using PySpark and SQL.
- Verify the contents of the table by displaying the first 5 rows.

#### 2. Write Data to Delta Lake:

- Append new transactions to the Delta table using PySpark.
- Example new transactions:

```
6,2024-09-06,C005,Keyboard,4,100
7,2024-09-07,C006,Mouse,10,20
```

#### 3. Update Data in Delta Lake:

- Update the `Price` of `Product = 'Laptop'` to `1300`.

- Use PySpark or SQL to perform the update and verify the results.

#### 4. Delete Data from Delta Lake:

- Delete all transactions where the `Quantity` is less than 3.
- Use both PySpark and SQL to perform this deletion.

#### 5. Merge Data into Delta Lake:

- Create a new set of data representing updates to the existing transactions. Merge the following new data into the Delta table:

```
TransactionID,TransactionDate,CustomerID,Product,Quantity,Price
1,2024-09-01,C001,Laptop,1,1250 -- Updated Price
8,2024-09-08,C007,Charger,2,30 -- New Transaction
```

- Use the Delta Lake **merge** operation to insert the new data and update the existing records.

---

### Exercise 3: Delta Lake - History, Time Travel, and Vacuum

#### Objective:

Understand how to use Delta Lake features such as versioning, time travel, and data cleanup with vacuum.

#### Tasks:

##### 1. View Delta Table History:

- Query the **history** of the Delta table to see all changes (inserts, updates, deletes) made in the previous exercises.
- Use both PySpark and SQL to view the history.

##### 2. Perform Time Travel:

- Retrieve the state of the Delta table as it was **5 versions ago**.
- Verify that the table reflects the data before some of the updates and deletions made earlier.
- Perform a query to get the transactions from a specific timestamp (e.g., just before an update).

##### 3. Vacuum the Delta Table:

- Clean up old data using the **VACUUM** command.
- Set a retention period of 7 days and vacuum the Delta table.
- Verify that old versions are removed, but the current table state is intact.

##### 4. Converting Parquet Files to Delta Files:

- Create a new Parquet-based table from the raw transactions CSV file.
- Convert this Parquet table to a Delta table using Delta Lake functionality.

---

### Exercise 4: Implementing Incremental Load Pattern using Delta Lake

#### Objective:

Learn how to implement incremental data loading with Delta Lake to avoid reprocessing old data.

**Tasks:**

**1. Set Up Initial Data:**

- Use the same transactions data from previous exercises, but load only transactions from the first three days ( 2024-09-01 to 2024-09-03 ) into the Delta table.

**2. Set Up Incremental Data:**

- Add a new set of transactions representing the next four days ( 2024-09-04 to 2024-09-07 ).
- Ensure that these transactions are loaded incrementally into the Delta table.

**3. Implement Incremental Load:**

- Create a pipeline that reads new transactions only (transactions after 2024-09-03 ) and appends them to the Delta table without overwriting existing data.
- Verify that the incremental load only processes new data and does not duplicate or overwrite existing records.

**4. Monitor Incremental Load:**

- Check the Delta Lake version history to ensure only the new transactions are added, and no old records are reprocessed.
-